

Bidirectional Elaboration of Dependently Typed Programs

Francisco Ferreira Brigitte Pientka

McGill University

{fferre8,bpientka}@cs.mcgill.ca

Abstract

Dependently typed programming languages allow programmers to express a rich set of invariants and verify them statically via type checking. To make programming with dependent types practical, dependently typed systems provide a compact language for programmers where one can omit some arguments, called implicit, which can be inferred. This source language is then usually elaborated into a core language where type checking and fundamental properties such as normalization are well understood. Unfortunately, this elaboration is rarely specified and in general is ill-understood. This makes it not only difficult for programmers to understand why a given program fails to type check, but also is one of the reasons that implementing dependently typed programming systems remains a black art known only to a few.

In this paper, we specify the design of a source language for a dependently typed programming language where we separate the language of programs from the language of types and terms occurring in types. Total functions in our language correspond directly to first-order inductive proofs over a specific index domain. We then give a bi-directional elaboration algorithm to translate source terms where implicit arguments can be omitted to a fully explicit core language and prove soundness of our elaboration. Our framework provides post-hoc explanation for elaboration found in the programming and proof environment, Beluga.

Categories and Subject Descriptors D.3.1 [Formal Definitions and Theory]: Syntax; F.3.3 [Studies of Program Constructs]: Type structure

Keywords dependent types, type reconstruction

1. Introduction

Dependently typed programming languages allow programmers to express a rich set of properties and statically verify them via type checking. To make programming with dependent types practical, these systems provide a source language where programmers can omit (implicit) arguments which can be reasonably easily inferred and elaborate the source language into a well-understood core language, an idea going back to Pollack [1990]. However, this elaboration is rarely specified formally for dependently typed languages which support recursion and pattern matching. For Agda, a full de-

pendently typed programming language based on Martin L f type theory, Norrel [2007, Chapter 3] describes a bi-directional type inference algorithm, but does not treat the elaboration of recursion and pattern matching. For the fully dependently typed language Idris, Brady [2013] describes the elaboration between source and target, but no theoretical properties such as soundness are established. A notable exception is Asperti et.al [2012] that describes a sound bi-directional elaboration algorithm for the Calculus of (Co)Inductive Constructions (CCIC) implemented in Matita.

In this paper, we concentrate on dependently typed programs which directly correspond to first-order logic proofs over a specific domain. More specifically, a proof by cases (and more generally by induction) corresponds to a (total) functional program with dependent pattern matching. We hence separate the language of programs from the language of our specific domain about which we reason. Our language is similar to indexed type systems (see [Zenger 1997; Xi and Pfenning 1999]); however, unlike these aforementioned systems, we allow pattern matching on index objects, i.e. we support case-analysis on objects in our domain. As a consequence, we cannot simply erase our implicit arguments and obtain a program which is simply typed.

Specifically, our source language is inspired by the Beluga language [Pientka 2008; Pientka and Dunfield 2010; Cave and Pientka 2012] where we specify formal systems in the logical framework LF [Harper et al. 1993] (our index language) and write proofs about LF objects as total recursive functions using pattern matching.

More generally, our language may be viewed as a smooth extension of simply typed languages, like Haskell or OCaml and we support nested pattern matching without having to either specify the return type or removing nested pattern matching during elaboration which is often the case in fully dependently typed systems such as Agda or Coq. Moreover, taking advantage of the separation between types and terms, it is easy to support effects, allow non-termination, partial functions, and polymorphism. All this while reaping some of the benefits of dependent types.

The main contribution of this paper is the design of a source language for dependently typed programs where we omit implicit arguments together with a sound bi-directional elaboration algorithm from the source language to a fully explicit core language. This language supports dependent pattern matching without requiring type invariant annotations, and dependently-typed case expressions can be nested as in simply-typed pattern matching. Throughout our development, we will keep the index language abstract and state abstractly our requirements such as decidability of equality and typing. There are many interesting choices of index languages. For example choosing arithmetic would lead to a DML [Xi 2007] style language ; choosing an authorization logic would let us manipulate authorization certificates (similar to *Aura* [Jia et al. 2008]); choosing LF style languages (like Contextual LF [Nanevski et al. 2008]) we obtain Beluga; choosing substructural variant of it like CLF [Watkins et al. 2002] we are in principle able to manipulate and work with substructural specifications.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

PPDP '14, September 8–10, 2014, Canterbury, UK.

Copyright is held by the owner/author(s). Publication rights licensed to ACM.

ACM 978-1-4503-2947-7/14/09...\$15.00.

http://dx.doi.org/10.1145/2643135.2643153

A central question when elaborating dependently typed languages is what arguments may the programmer omit. In dependently-typed systems such as Agda or Coq, the programmer declares constants of a given (closed) type and labels arguments that can be freely omitted when subsequently using the constant. Both, Coq and Agda, give the user the possibility to locally override the implicit arguments and provide instantiations explicitly.

In contrast, we follow here a simple, lightweight recipe which comes from the implementation of the logical framework *Elf* [Pfenning 1989] and its successor *Twelf* [Pfenning and Schürmann 1999]: *programmers may leave some index variables free when declaring a constant of a given type; elaboration of the type will abstract over these free variables at the outside; when subsequently using this constant, the user must omit passing arguments for those index variables which were left free in the original declaration.* Following this recipe, elaboration of terms and types in the logical framework has been described in Pientka [2013]. Here, we will consider a dependently typed functional programming language which supports pattern matching on index objects.

The key challenge in elaborating recursive programs which support case-analysis is that pattern matching in the dependently typed setting refines index arguments and hence refines types. In contrast to systems such as Coq and Agda, where we must annotate case-expressions with an invariant, i.e. the type of the scrutinee, and the return type, our source language does not require such annotations. Instead we will infer the type of the scrutinee and for each branch, we infer the type of the pattern and compute how the pattern refines the type of the scrutinee. This makes our source language lightweight and closer in spirit to simply-typed functional languages. Our elaboration of source expressions to target expressions is type-directed, inferring omitted arguments and producing a closed well-typed target program. Finally, we prove soundness of our elaboration, i.e. if elaboration succeeds our resulting program type checks in our core language. Our framework provides post-hoc explanation for elaboration found in the programming and proof environment, Beluga [Pientka 2008; Cave and Pientka 2012], where we use as the index domain terms specified in the logical framework LF [Harper et al. 1993].

The paper is organized as follows: We first give the grammar of our source language. Showing example programs, we explain informally what elaboration does. We then revisit our core language, describe the elaboration algorithm formally and prove soundness. We conclude with a discussion of related and future work.

2. Source language

We consider here a dependently typed language where types are indexed by terms from an index domain. Our language is similar to Beluga [Pientka and Dunfield 2010], a dependently typed programming environment where we specify formal systems in the logical framework LF and we can embed LF objects into computation-level types and computation-level programs which analyze and pattern match on LF objects. However, in our description, as in for example Cave and Pientka [2012], we will keep the index domain abstract, but only assume that equality in the index domain is decidable and unification algorithms exist. This will allow us to focus on the essential challenges when elaborating a dependently typed language in the presence of pattern matching.

We describe the *source language* that allows programmers to omit some arguments in Fig. 1. As a convention we will use lowercase c to refer to index level objects, lowercase u for index level types, and upper case letters X, Y for index-variables. Index objects can be embedded into computation expressions by using a box modality written as $[c]$. Our language supports functions ($\mathbf{fn} x \Rightarrow e$), dependent functions ($\lambda X \Rightarrow e$), function application ($e_1 e_2$), dependent function application ($e_1 [c]$), and case-expressions. We also

Kinds	$k ::= \mathbf{ctype} \mid \{X:u\} k$
Atomic Types	$p ::= \mathbf{a} [\vec{c}]$
Types	$t ::= p \mid [u] \mid \{X:u\} t \mid t_1 \rightarrow t_2$
Expressions	$e ::= \mathbf{fn} x \Rightarrow e \mid \lambda X \Rightarrow e \mid x \mid \mathbf{c} \mid [c] \mid e_1 e_2 \mid e_1 [c] \mid e _ \mid \mathbf{case} e \mathbf{of} \vec{b} \mid e:t$
Branches	$\vec{b} ::= b \mid (b \mid \vec{b})$
Branch	$b ::= pat \mapsto e$
Pattern	$pat ::= x \mid [c] \mid \mathbf{c} \vec{pat} \mid pat:t$
Declarations	$d ::= \mathbf{rec} f:t = e \mid \mathbf{c}:t \mid \mathbf{a}:k$

Figure 1. Grammar of Source Language

support writing underscore ($_$) instead of providing explicitly the index argument in a dependent function application ($e _$). Note that we are overloading syntax: we write $e [c]$ to describe the application of the expression e of type $[u] \rightarrow t$ to the expression $[c]$; we also write $e [c]$ to describe the dependent application of the expression e of type $\{X:u\}t$ to the (unboxed) index object c . This ambiguity can be easily resolved using type information. Note that in our language the dependent function type and the non-dependent function type do not collapse, since we can only quantify over objects of our specific domain instead of arbitrary propositions (types).

We may write type annotations anywhere in the program ($e:t$ and in patterns $pat:t$); type annotations are particularly useful to make explicit the type of a sub-expression and name index variables occurring in the type. This allows us to resurrect index variables which are kept implicit. In patterns, type annotations are useful since they provide hints to type elaboration regarding the type of pattern variables.

A program signature Σ consists of kind declarations ($\mathbf{a}:k$), type declarations ($\mathbf{c}:t$) and declaration of recursive functions ($\mathbf{rec} f:t = e$). This can be extended to allow mutual recursive functions in a straightforward way.

One may think of our source language as the language obtained after parsing where for example let-expressions have been translated into case-expression with one branch.

Types for computations include non-dependent function types ($t_1 \rightarrow t_2$) and dependent function types ($\{X:u\}t$); we can also embed index types into computation types via $[u]$ and indexed computation-level types by an index domain written as $\mathbf{a} [\vec{c}]$. We also include the grammar for computation-level kinds which emphasizes that computation-level types can only be indexed by terms from an index domain u . We write \mathbf{ctype} (i.e. computation-level type) for the base kind, since we will use \mathbf{type} for kinds of the index domain.

We note that we do only support one form of dependent function type $\{X:u\}t$; the source language does not provide any means for programmers to mark a given dependently typed variable as implicit as for example in Agda. Instead, we will allow programmers to leave some index variables occurring in computation-level types free; elaboration will then infer their types and abstract over them explicitly at the outside. The programmer must subsequently omit providing instantiation for those “free” variables. We will explain this idea more concretely below.

2.1 Well-formed source expressions

Before elaborating source expressions, we state when a given source expression is accepted as a well-formed expression. In particular, it will highlight that free index variables are only allowed in declarations when specifying kinds and declaring the type of constants and recursive functions. We use δ to describe the list of index variables and γ the list of program variables. We rely on two

$\boxed{\vdash d \text{ wf}}$ Declaration d is well-formed

$$\frac{\cdot; f \vdash e \text{ wf} \quad \cdot \vdash_f t \text{ wf}}{\vdash \text{rec } f:t = e \text{ wf}} \text{ wf-rec} \quad \frac{\cdot \vdash_f t \text{ wf}}{\vdash c:t \text{ wf}} \text{ wf-types} \quad \frac{\cdot \vdash_f k \text{ wf}}{\vdash a:k \text{ wf}} \text{ wf-kinds}$$

$\boxed{\delta; \gamma \vdash_N e \text{ wf}}$ Normal expression e is well-formed in context δ and γ

$$\frac{\delta; \gamma, x \vdash_N e \text{ wf}}{\delta; \gamma \vdash_N \text{fn } x \Rightarrow e \text{ wf}} \text{ wf-fn} \quad \frac{\delta, X; \gamma \vdash_N e \text{ wf}}{\delta; \gamma \vdash_N \lambda X \Rightarrow e \text{ wf}} \text{ wf-mlam} \quad \frac{\delta \vdash c \text{ wf}}{\delta; \gamma \vdash_N [c] \text{ wf}} \text{ wf-box}$$

$$\frac{\delta; \gamma \vdash_n e \text{ wf} \quad \text{for all } b_n \text{ in } \vec{b}. \delta; \gamma \vdash b_n \text{ wf}}{\delta; \gamma \vdash_N \text{case } e \text{ of } \vec{b} \text{ wf}} \text{ wf-case} \quad \frac{\delta; \gamma \vdash_n e \text{ wf}}{\delta; \gamma \vdash_N e \text{ wf}} \text{ wf-neu}$$

$\boxed{\delta; \gamma \vdash_n e \text{ wf}}$ Neutral expression e is well-formed in context δ and γ

$$\frac{\delta; \gamma \vdash_N e \text{ wf} \quad \delta \vdash t \text{ wf}}{\delta; \gamma \vdash_n e:t \text{ wf}} \text{ wf-ann} \quad \frac{\delta; \gamma \vdash_n e_1 \text{ wf} \quad \delta; \gamma \vdash_N e_2 \text{ wf}}{\delta; \gamma \vdash_n e_1 e_2 \text{ wf}} \text{ wf-app} \quad \frac{\delta; \gamma \vdash_n e \text{ wf} \quad \delta \vdash c \text{ wf}}{\delta; \gamma \vdash_n e[c] \text{ wf}} \text{ wf-app-explicit}$$

$$\frac{\delta; \gamma \vdash_n e_1 \text{ wf}}{\delta; \gamma \vdash_n e_1 - \text{wf}} \text{ wf-apph} \quad \frac{x \in \gamma}{\delta; \gamma \vdash_n x \text{ wf}} \text{ wf-var}$$

$\boxed{\delta; \gamma \vdash \text{pat} \mapsto e \text{ wf}}$ Branch is well-formed in δ and γ

$$\frac{\delta'; \gamma' \vdash \text{pat} \text{ wf} \quad \delta, \delta'; \gamma, \gamma' \vdash_n e \text{ wf}}{\delta; \gamma \vdash \text{pat} \mapsto e \text{ wf}} \text{ wf-branch}$$

$\boxed{\delta; \gamma \vdash \text{pat} \text{ wf}}$ Pattern pat is well-formed in a context δ for index variables and a context γ for pattern variables

$$\frac{}{\delta; x \vdash x \text{ wf}} \text{ wf-p-var} \quad \frac{\delta \vdash c \text{ wf}}{\delta; \cdot \vdash [c] \text{ wf}} \text{ wf-p-i} \quad \frac{\text{for all } p_i \text{ in } \vec{Pat}. \delta; \gamma_i \vdash p_i \text{ wf}}{\delta; \gamma_1, \dots, \gamma_n \vdash \text{c } \vec{Pat} \text{ wf}} \text{ wf-p-con} \quad \frac{\delta; \gamma \vdash \text{pat} \text{ wf} \quad \delta \vdash t \text{ wf}}{\delta; \gamma \vdash \text{pat}:t \text{ wf}} \text{ wf-p-ann}$$

Figure 2. Well-formed source expressions

judgments from the index language:

- $\delta \vdash c \text{ wf}$ Index object c is well formed and closed with respect to δ
- $\delta \vdash_f c \text{ wf}$ Index object c is well formed with respect to δ and may contain free index variables

We describe declaratively the well-formedness of declarations and source expressions in Fig. 2. The distinction between normal and neutral expressions forces a type annotation where a non-normal program would occur. The normal vs. neutral term distinction is motivated by the bidirectional type-checker presented in Section 3.1. For brevity, we omit the full definition of well-formedness for kinds and types which is given in the appendix.

In branches, pattern variables from γ must occur linearly while we put no such requirement on variables from our index language listed in δ . The judgement for well formed patterns synthesizes contexts δ and γ that contain all the variables bound in the pattern (this presentation is declarative, but algorithmically the two contexts result from the well-formed judgement). Notice that the rules wp-p-i and wp-p-con look similar but they operate on different syntactic categories and refer to the judgement for well-formed index terms provided by the index level language. They differ in that the one for patterns synthesizes the δ context that contains the meta-variables bound in the pattern.

2.2 Some example programs

We next illustrate writing programs in our language and explain the main ideas behind elaboration.

2.2.1 Translating untyped terms to intrinsically typed terms

We implement a program to translate a simple language with numbers, booleans and some primitive operations to its typed counterpart. This illustrates declaring an index domain, using index computation-level types, and explaining the use and need to pattern match on index objects. We first define the untyped version of our language using the recursive datatype UTm . Note the use of the keyword **ctype** to define a computation-level recursive data-type.

```
datatype UTm : ctype =
| UNum : Nat → UTm
| UPlus : UTm → UTm → UTm
| UTrue : UTm
| UFalse : UTm
| UNot : UTm → UTm
| UIf : UTm → UTm → UTm → UTm;
```

Terms can be of type `nat` for numbers or `bool` for booleans. Our goal is to define our language of typed terms using a computation-level type family Tm which is indexed by objects `nat` and `bool` which are constructors of our index type `tp`. Note that `tp` is declared as having the kind **type** which implies that this type lives at the index level and that we will be able to use it as an index for computation-level type families.

```
datatype tp : type =
| nat : tp
| bool : tp;
```

Using indexed families we can now define the type Tm that specifies only type correct terms of the language, by indexing terms by their type using the index level type tp .

```
datatype Tm : [tp] → ctype =
| Num      : Nat      → Tm [nat]
| Plus     : Tm [nat] → Tm [nat] → Tm [nat]
| True     : Tm [bool]
| False    : Tm [bool]
| Not      : Tm [bool] → Tm [bool]
| If       : Tm [bool] → Tm [T] → Tm [T] → Tm [T];
```

When the Tm family is elaborated, the free variable T in the If constructor will be abstracted over by an implicit Π -type, as in the Twelf [Pfenning and Schürmann 1999] tradition. Because T was left free by the programmer, the elaboration will add an implicit quantifier; when we use the constant If , we now must omit passing an instantiation for T . For example, we must write $(If\ True\ (Num\ 3)\ (Num\ 4))$ and elaboration will infer that T must be nat .

One might ask how we can provide the type explicitly - this is possible indirectly by providing type annotations. For example, $If\ e\ (e1 : Tm[nat])\ e2$ will fix the type of $e1$ to be $Tm\ [nat]$.

Our goal is to write a program to translate an untyped term UTm to its corresponding typed representation. Because this operation might fail for ill-typed UTm terms we need an option type to reflect the possibility of failure.

```
datatype TmOpt : ctype =
| None : TmOpt
| Some : {T : tp} Tm [T] → TmOpt;
```

A value of type $TmOpt$ will either be empty (i.e. $None$) or some term of type T . We chose to make T explicit here by quantifying over it explicitly using the curly braces. When returning a Tm term, the program must now provide the instantiation of T in addition to the actual term.

So far we have declared types and constructors for our language. These declarations will be available in a global signature. The next step is to declare a function that will take untyped terms into typed terms if possible. Notice that for the function to be type correct it has to respect the specification provided in the declaration of the type Tm . We only show a few interesting cases below.

```
rec tc : UTm → TmOpt = fn e ⇒ case e of ...
| UNum n ⇒ Some [nat] (Num n)
| UNot e ⇒ (case tc e of
| Some [bool] x ⇒ Some [bool] (Not x)
| other ⇒ None)
| UIf c e1 e2 ⇒ (case tc c of
| Some [bool] c' ⇒ (case (tc e1 , tc e2) of
| (Some [T] e1' , Some [T] e2') ⇒
Some [T] (If c' e1' e2')
| other ⇒ None )
| other ⇒ None )
;
```

In the tc function the cases for numbers, plus, true and false are completely straightforward. The case for negation (i.e. constructor $UNot$) is interesting because we need to pattern match on the result of type-checking the sub-expression e to match its type to $bool$ otherwise we cannot construct the intrinsically typed term, i.e. the constructor Not requires a boolean term, this requires matching on index level terms. Additionally the case for UIf is also interesting because we not only need a boolean condition but we also need to have both branches of the UIf term to be of the same type. Again we use pattern matching on the indices to verify that the condition is of type $bool$ but notably we use non-linear pattern matching to ensure that the type of the branches coincides. We note that If has an implicit argument (the type T) which will be inferred during elaboration.

In the definition of type $TmOpt$ we chose to explicitly quantify over T , however another option would have been to leave it implicit. When pattern matching on $Some\ e$, we would need to resurrect the type of the argument e to be able to inspect it and check whether it has the appropriate type. We can employ type annotations, as shown in the code below, to constrain the type of e .

```
| UIf c e1 e2 ⇒ (case tc c of
| Some (c' : Tm [bool]) ⇒ (case (tc e1 , tc e2) of
| (Some (e1' : Tm [T]) , Some (e2' : Tm [T])) ⇒
Some (If c' e1' e2')
| other ⇒ None)
| other ⇒ None)
```

In this first example there is not much to elaborate. The missing argument in If and the types of variables in patterns are all that need elaboration.

2.2.2 Type-preserving evaluation

Our previous program used dependent types sparingly; most notably there were no dependent types in the type declaration given to the function tc . We now discuss the implementation of an evaluator, which evaluates type correct programs to values of the same type, to highlight writing dependently typed functions. Because we need to preserve the type information, we index the values by their types in the following manner:

```
datatype Val : [tp] → ctype =
| VNum : Nat → Val [nat]
| VTrue : Val [bool]
| VFalse : Val [bool];
```

We define a type preserving evaluator below; again, we only show some interesting cases.

```
rec eval : Tm [T] → Val [T] = fn e ⇒ case e of ...
| Num n ⇒ VNum n
| Plus e1 e2 ⇒ (case (eval e1 , eval e2) of
| (VNum x , VNum y) ⇒ VNum (add x y))
| Not e ⇒ (case eval e of
| VTrue ⇒ VFalse
| VFalse ⇒ VTrue)
| If e e1 e2 ⇒ (case eval e of
| VTrue ⇒ eval e1
| VFalse ⇒ eval e2)
;
```

First, we specify the type of the evaluation function as $Tm[T] \rightarrow Val[T]$ where T remains free. As a consequence, elaboration will infer its type and abstract over T at the outside, T is an implicit parameter (as it is introduced by Π^i). We now elaborate the body of the function against $\Pi^i T : tp.\ Tm\ [T] \rightarrow Val\ [T]$. It will first need to introduce the appropriate dependent function abstraction in the program before we introduce the non-dependent function $fn\ x \Rightarrow e$. Moreover, we need to infer omitted arguments in the pattern in addition to inferring the type of pattern variables in the If case. Since T was left free in the type given to $eval$, we must also infer the omitted argument in the recursive calls to $eval$. Finally, we need to keep track of refinements the pattern match induces: our scrutinee has type $Tm\ [T]$; pattern matching against $Plus\ e1\ e2$ which has type $Tm\ [nat]$ refines T to nat .

2.2.3 A certifying evaluator

So far in our examples, we have used a simply typed index language. We used our index language to specify natural numbers, booleans, and a tagged enumeration that contained labels for the $bool$ and nat types. In this example we go one step further, and use a dependently typed specification, in fact we use LF as our index level language as used in Beluga [Cave and Pientka 2012]. Using

LF at the index language we specify the simply-typed lambda calculus and its operational semantics in Figure 3. Using these specifications we write a recursive function that returns the value of the program together with a derivation tree that shows how the value was computed. This example requires dependent types at the index level and consequently the elaboration of functions that manipulate these specifications has to be more powerful.

As in the previous example, we define the types of terms of our language using the index level language. As opposed to the type preserving evaluator, in this case we define our intrinsically typed terms also using the index level language (which will be LF for this example). We take advantage of LF to represent binders in the lambda terms, and use dependent types to represent well-typed terms only.

```
datatype tp : type =
| unit : tp
| arr : tp → tp → tp
;

datatype term : tp → type =
| one : term unit
| lam : (term A → term B) → term (arr A B)
| app : term (arr A B) → term A → term B
;
```

These datatypes represent an encoding of well typed terms of a simply-typed lambda calculus with `unit` as a base type. Using LF we can also describe what constitutes a value and a big-step operational semantics. We use the standard technique in LF to represent binders with the function space (usually called Higher Order Abstract Syntax, HOAS [Pfenning and Elliott 1988]) and type families to only represent well-typed terms, thus this representation combines the syntax for terms with the typing judgement from Figure 3.

```
datatype value : tp → type =
| v-one : value unit
| v-lam : (term A → term B) → value (arr A B)
;

datatype big-step : term T → value T → type =
| e-one : big-step one v-one
| e-lam : big-step (lam M) (v-lam M)
| e-app : big-step M (v-lam M') →
          big-step (M' N) N' →
          big-step (app M N) N'
;
```

The value type simply states that one and lambda terms are values, and the type `big-step` encodes the operational semantics where each constructor corresponds to one of the rules in Figure 3. The constructors `e-one` and `e-lam` simply state that `one` and `lambda` step to themselves. On the other hand rule `e-app` requires that in an application, the first term evaluates to a lambda expression (which is always the case as the terms are intrinsically well typed) and then it performs the substitution and continues evaluating the term to a value. Note how the substitution is performed by an application as we reuse the LF function space for binders as typically done with HOAS.

In this paper we discuss how to elaborate recursive programs, however because this example uses dependently typed specifications at the index level, these specifications will be elaborated following the elaboration described in Pientka [2013].

To implement a certifying evaluator we want the `eval` function to return a value and a derivation tree that shows how we computed this value. We encode this fact in the `Cert` data-type that encodes an existential or dependent pair that combines a value with a derivation tree.

Types $T ::= \top \mid T_1 \rightarrow T_2$
Terms $M, N ::= () \mid x \mid \lambda x:T.M \mid MN$
Context $\Gamma ::= \cdot \mid \Gamma, x:T$
Values $V ::= \top \mid \lambda x:T.M$

$\boxed{\Gamma \vdash M:T}$ Term M has type T in context Γ

$$\frac{}{\Gamma \vdash () : \top} \quad \frac{x:T \in \Gamma}{\Gamma \vdash x:T} \quad \frac{\Gamma, x:T_1 \vdash M:T_2}{\Gamma \vdash \lambda x:T_1.M:T_2}$$

$$\frac{\Gamma \vdash M:T_1 \rightarrow T \quad \Gamma \vdash N:T_1}{\Gamma \vdash MN:T}$$

$\boxed{M \Downarrow V}$ Term M steps to value V

$$\frac{}{\top \Downarrow \top} \quad \frac{}{\lambda x:T.M \Downarrow \lambda x:T.M}$$

$$\frac{M \Downarrow \lambda x:T.M' \quad [N/x]M' \Downarrow N'}{MN \Downarrow N'}$$

Figure 3. Example: A simply-typed lambda calculus

```
datatype Cert : [term T] → ctype =
| Ex : {N: [value T]}[big-step M N] → Cert [M]
;
```

In the `Ex` constructor we have chosen to explicitly quantify over `N`, the value of the evaluation, and left the starting term `M` implicit. However another option would have been to leave both implicit, and use type annotations when pattern matching to have access to both the term and its value.

Finally the evaluation function simply takes a term and returns a certificate that contains the value the terms evaluates to, and the derivation tree that led to that value.

```
rec eval : {M : [term T]} Cert [M] =
λ M ⇒ case [M] of
| [one] ⇒ Ex [v-one] [e-one]
| [lam M] ⇒ Ex [v-lam M] [e-lam]
| [app M N] ⇒
  let Ex [v-lam M'] [D] = eval [M] in
  let Ex [N'] [D'] = eval [M' N] in
  Ex [N'] [e-app D D']
;
```

Elaboration of `eval` starts by the type annotation. Inferring the type of variable `T` and abstracting over it, resulting in:

$\Pi^T T : [tp]. \{M : [term T] \rightarrow [value T]\}$. The elaboration proceeds with the body, abstracting over the inferred dependent argument with $\lambda T \Rightarrow \dots$. When elaborating the case expression, the patterns in the index language will need elaboration. In this work we assume that each index language comes equipped with an appropriate notion of elaboration (described in [Pientka 2013] for this example). For example, index level elaboration will abstract over free variables in constructors and the pattern for lambda terms becomes $[\text{lam } A \ B \ (\lambda x. \ M \ x)]$ when the types for parameters and body are added (and the body of the lambda is η -expanded). Additionally, in order to keep the core language as lean as possible we desugar `let` expressions into `case` expressions. For example, in the certifying evaluator, the following code from `eval`:

```
let Ex [v-lam M'] [D] = eval [M] in
let Ex [N'] [D'] = eval [M' N] in
Ex [N'] [e-app D D']
```

Kinds	$K ::= \Pi^e X:U. K \mid \Pi^i X:U. K \mid \mathbf{ctype}$
Atomic Types	$P ::= \mathbf{a} \vec{C}$
Types	$T ::= \Pi^e X:U. T \mid \Pi^i X:U. T \mid P \mid [U] \mid T_1 \rightarrow T_2$
Expressions	$E ::= \mathbf{fn} x \Rightarrow E \mid \lambda X \Rightarrow E \mid E_1 E_2 \mid E_1 [C] \mid [C] \mid \mathbf{case} E \mathbf{of} \vec{B} \mid x \mid E:T \mid \mathbf{c}$
Branches	$\vec{B} ::= B \mid (B \mid \vec{B})$
Branch	$B ::= \Pi \Delta; \Gamma. Pat: \theta \mapsto E$
Pattern	$Pat ::= x \mid [C] \mid \mathbf{c} \vec{Pat}$
Declarations	$D ::= \mathbf{c}:T \mid \mathbf{a}:K \mid \mathbf{rec} f:T = E$
Context	$\Gamma ::= \cdot \mid \Gamma, x:T$
Index-Var-Context	$\Delta ::= \cdot \mid \Delta, X:U$
Refinement	$\theta ::= \cdot \mid \theta, C/X \mid \theta, X/X$

Figure 4. Target language

is desugared into:

```
(case eval [M] of | Ex [v-lam M'] [D] =>
  (case eval [M' N] of
    | Ex [N'] [D'] => Ex [N'] [e-app D D'])))
```

We will come back to this example and discuss the fully elaborated program in the next section.

3. Target language

The *target language* is similar to the computational language described in Cave and Pientka [2012] which has a well developed meta-theory including descriptions of coverage [Dunfield and Pientka 2009] and termination [Pientka et al. 2014]. The target language (see Fig. 4), which is similar to our source language, is indexed by fully explicit terms of the index level language; we use C for fully explicit index level objects, and U for elaborated index types; index-variables occurring in the target language will be represented by capital letters such as X, Y . Moreover, we rely on a substitution which replaces index variables X with index objects. The main difference between the source and target language is in the description of branches. In each branch, we make the type of the pattern variables (see context Γ) and variables occurring in index objects (see context Δ) explicit. We associate each pattern with a refinement substitution θ which specifies how the given pattern refines the type of the scrutinee.

3.1 Typing of target language

The typing rules for our core language are given in Fig. 5. We again omit the rules for types and kind which are given in Cave and Pientka [2012].

We use a bidirectional type system [Pierce and Turner 2000] for the target language which is similar to the one in Cave and Pientka [2012] but we simplify the presentation by omitting recursive types. Instead we assume that constructors together with their types are declared in a signature Σ . We choose a bi-directional type-checkers because it minimizes the need for annotations by propagating known typing information in the check phase (judgement $\Delta; \Gamma \vdash E \Leftarrow T$ in the rules) and inferring the types when it is possible in the synthesis phase (judgement $\Delta; \Gamma \vdash E \Rightarrow T$).

We rely on the fact that our index domain comes with rules which check that a given index object is well-typed. This is described by the judgment: $\Delta \vdash C : U$.

The introductions, functions $\mathbf{fn} x \Rightarrow e$ and dependent functions $\lambda x \Rightarrow e$, check against their respective type. Dependent functions check against both $\Pi^e X:U. T$ and $\Pi^i X:U. T$ where types are annotated with e for explicit quantification and i for implicit quantification filled in by elaboration. Their corresponding eliminations, application $E_1 E_2$ and dependent application $E [C]$, synthesize their type. We rely in this rule on the index-level substitution operation and we assume that it is defined in such a way that normal forms are preserved¹.

To type-check a case-expressions $\mathbf{case} E \mathbf{of} \vec{B}$ against T , we synthesize a type S for E and then check each branch against $S \rightarrow T$. A branch $\Pi \Delta'; \Gamma'. Pat: \theta \mapsto E$ checks against $S \rightarrow T$, if: 1) θ is a refinement substitution mapping all index variables declared in Δ to a new context Δ' , 2) the pattern Pat is compatible with the type S of the scrutinee, i.e. Pat has type $[\theta]S$, and the body E checks against $[\theta]T$ in the index context Δ' and the program context $[\theta]\Gamma, \Gamma_i$. Note that the refinement substitution effectively performs a context shift.

We present the typing rules for patterns in spine format which will simplify our elaboration and inferring types for pattern variables. We start checking a pattern against a given type and check index objects and variables against the expected type. If we encounter $\mathbf{c} \vec{Pat}$ we look up the type T of the constant \mathbf{c} in the signature and continue to check the spine \vec{Pat} against T with the expected return type S . Pattern spine typing succeeds if all patterns in the spine \vec{Pat} have the corresponding type in T and yields the return type S .

3.2 Elaborated examples

In Section 2.2.3 we give an evaluator for a simply typed lambda calculus that returns the result of the evaluation together with the derivation tree needed to compute the value. The elaborated version of function `eval` is:

```
rec eval :  $\Pi^i. T:tp. \{M : [\text{term } T]\} \text{Cert } [T] [M] =$ 
 $\lambda T \Rightarrow \lambda M \Rightarrow \mathbf{case} [M] \mathbf{of}$ 
|  $\dots [\text{one}] : \text{unit}/T \Rightarrow \text{Ex } [\text{unit}] [\text{v-one}] [\text{e-one}]$ 
|  $T1:[tp], T2:[tp], M : [x:T1 \vdash \text{term } T2];.$ 
   $[\text{lam } T1 T2 (\lambda x. M x)] : \text{arr } T1 T2 / T \Rightarrow$ 
   $\text{Ex } [\text{arr } T1 T2]$ 
   $[\text{v-lam } T1 T2 (\lambda x. M x)]$ 
   $[\text{e-lam } (\lambda x. M x)]$ 
|  $T1:[tp], T2:[tp],$ 
 $M:[\text{term } (\text{arr } T1 T2)], N:[\text{term } T1];.$ 
   $[\text{app } T1 T2 M N] : T2/T \Rightarrow$ 
   $(\mathbf{case} \text{eval } [\text{arr } T1 T2] [M] \mathbf{of}$ 
  |  $T1:[tp], T2:[tp], M':[x:T1 \vdash \text{term } T2],$ 
     $D : [\text{big-step } (\text{arr } T1 T2) M'$ 
       $(\text{v-lam } (\lambda x. M x))];.$ 
     $\text{Ex } [\text{arr } T1 T2] [\text{v-lam } (\text{arr } T1 T2) M'] [D] : \dots \Rightarrow$ 
     $(\mathbf{case} \text{eval } [T2] [M' N] \mathbf{of}$ 
    |  $T2:[tp], N':[\text{val } T2],$ 
       $D':[\text{big-step } T2 (M' N) N'];.$ 
       $\text{Ex } [T2] [N'] [D'] : \dots \Rightarrow$ 
       $\text{Ex } [T2] [N']$ 
       $[\text{e-app } M M' N N' D D'])))$ 
;
```

To elaborate a recursive declaration we start by reconstructing the type annotation given to the recursive function. In this case the user left the variable T free which becomes an implicit argument and we abstract over this variable with $\Pi^i. T:Tp$ marking it implicit. Notice however how the user explicitly quantified over M this means that callers of `eval` have to provide the term M while parameter T will be omitted and inferred at each calling point. Next,

¹ In Beluga, this is for example achieved by relying on hereditary substitutions [Cave and Pientka 2012].

$$\boxed{\vdash D \text{ wf}} \quad \text{Target declaration } D \text{ is well-formed}$$

$$\frac{\cdot \vdash T \Leftarrow \mathbf{ctype} \quad \cdot; f:T \vdash E \Leftarrow T}{\vdash \mathbf{rec} f:T = E \text{ wf}} \quad \mathbf{t-rec} \quad \frac{\cdot \vdash T \Leftarrow \mathbf{ctype}}{\vdash \mathbf{c}:T \text{ wf}} \quad \mathbf{t-type} \quad \frac{\cdot \vdash K \Leftarrow \mathbf{kind}}{\vdash \mathbf{a}:K \text{ wf}} \quad \mathbf{t-kind}$$

$$\boxed{\Delta; \Gamma \vdash E \Rightarrow T} \quad E \text{ synthesizes type } T$$

$$\frac{\Delta; \Gamma \vdash E_1 \Rightarrow S \rightarrow T \quad \Delta; \Gamma \vdash E_2 \Leftarrow S}{\Delta; \Gamma \vdash E_1 E_2 \Rightarrow T} \quad \mathbf{t-app} \quad \frac{\Delta; \Gamma \vdash E \Rightarrow \Pi^* X:U. T \quad * = \{i, e\} \quad \Delta \vdash C \Leftarrow U}{\Delta; \Gamma \vdash E [C] \Rightarrow [C/X]T} \quad \mathbf{t-app-index}$$

$$\frac{\Gamma(x) = T}{\Delta; \Gamma \vdash x \Rightarrow T} \quad \mathbf{t-var} \quad \frac{\Delta; \Gamma \vdash E \Leftarrow T}{\Delta; \Gamma \vdash E:T \Rightarrow T} \quad \mathbf{t-ann}$$

$$\boxed{\Delta; \Gamma \vdash E \Leftarrow T} \quad E \text{ type checks against type } T$$

$$\frac{\Delta; \Gamma \vdash E \Rightarrow T}{\Delta; \Gamma \vdash E \Leftarrow T} \quad \mathbf{t-syn} \quad \frac{\Delta; \Gamma, x:T_1 \vdash E \Leftarrow T_2}{\Delta; \Gamma \vdash (\mathbf{fn} x \Rightarrow E) \Leftarrow T_1 \rightarrow T_2} \quad \mathbf{t-fn}$$

$$\frac{\Delta, X:U; \Gamma \vdash E \Leftarrow T \quad * = \{i, e\}}{\Delta; \Gamma \vdash (\lambda X \Rightarrow E) \Leftarrow \Pi^* X:U. T} \quad \mathbf{t-mlam} \quad \frac{\Delta; \Gamma \vdash E \Rightarrow S \quad \Delta; \Gamma \vdash \vec{B} \Leftarrow S \rightarrow T}{\Delta; \Gamma \vdash \mathbf{case} E \text{ of } \vec{B} \Leftarrow T} \quad \mathbf{t-case}$$

$$\boxed{\Delta; \Gamma \vdash \Pi \Delta'; \Gamma'. Pat:\theta \mapsto E \Leftarrow T} \quad \text{Branch } B = \Pi \Delta'; \Gamma'. Pat:\theta \mapsto E \text{ checks against type } T$$

$$\frac{\Delta' \vdash \theta:\Delta \quad \Delta'; \Gamma' \vdash Pat \Leftarrow [\theta]S \quad \Delta'; [\theta]\Gamma, \Gamma' \vdash E \Leftarrow [\theta]T}{\Delta; \Gamma \vdash \Pi \Delta'; \Gamma'. Pat:\theta \mapsto E \Leftarrow S \rightarrow T} \quad \mathbf{t-branch}$$

$$\boxed{\Delta; \Gamma \vdash Pat \Leftarrow T} \quad \text{Pattern } Pat \text{ checks against } T$$

$$\frac{\Delta \vdash C \Leftarrow U}{\Delta; \Gamma \vdash [C] \Leftarrow [U]} \quad \mathbf{t-pindex} \quad \frac{\Gamma(x) = T}{\Delta; \Gamma \vdash x \Leftarrow T} \quad \mathbf{t-pvar} \quad \frac{\Sigma(\mathbf{c}) = T \quad \Delta; \Gamma \vdash \vec{Pat} \Leftarrow T \rangle S}{\Delta; \Gamma \vdash \mathbf{c} \vec{Pat} \Leftarrow S} \quad \mathbf{t-pcon}$$

$$\boxed{\Delta; \Gamma \vdash \vec{Pat} \Leftarrow T \rangle S} \quad \text{Pattern spine } \vec{Pat} \text{ checks against } T \text{ and has result type } S$$

$$\frac{\Delta \vdash C \Leftarrow U \quad \Delta; \Gamma \vdash \vec{Pat} \Leftarrow [C/X]T \rangle S}{\Delta; \Gamma \vdash [C] \vec{Pat} \Leftarrow \Pi^e X:U. T \rangle S} \quad \mathbf{t-spi} \quad \frac{\Delta; \Gamma \vdash Pat \Leftarrow T_1 \quad \Delta; \Gamma \vdash \vec{Pat} \Leftarrow T_2 \rangle S}{\Delta; \Gamma \vdash Pat \vec{Pat} \Leftarrow T_1 \rightarrow T_2 \rangle S} \quad \mathbf{t-sarr} \quad \frac{}{\Delta; \Gamma \vdash \cdot \Leftarrow S \rangle S} \quad \mathbf{t-snil}$$

Figure 5. Typing of computational expressions

we elaborate the function body given the fully elaborated type. We therefore add the corresponding abstraction $\lambda T \Rightarrow$ for the implicit parameter.

Elaboration proceeds recursively on the term. We reconstruct the case-expression, considering the scrutinee $[M]$ and we infer its type as $[\mathbf{term} \ T]$. We elaborate the branches next. Recall that each branch in the source language consists of a pattern and a body. Moreover, the body can refer to any variable in the pattern or variables introduced in outer patterns. However, in the target language branches abstract over the context $\Delta; \Gamma$ and add a refinement substitution θ . The body of the branch refers to variables declared in the branch contexts only. In each branch, we list explicitly the index variables and pattern variables. For example in the branch for $[\mathbf{lam} \ M]$ we added T_1 and T_2 to the index context Δ of the branch, index-level reconstruction adds these variables to the pattern. The refinement substitution moves terms from the outer context to the branch context, refining the appropriate index variables as expressed by the pattern. For example in this branch, the substitution refines the type $[T]$ to $[\mathbf{arr} \ T_1 \ T_2]$. And in the $[\mathbf{one}]$ branch it refines the type $[T]$ to $[\mathbf{unit}]$.

As we mentioned before, elaboration adds an implicit parameter to the type of function \mathbf{eval} , and the user is not allowed to directly supply an instantiation for it. Implicit parameters have to be inferred by elaboration. In the recursive calls to \mathbf{eval} , we add the parameter that represents the type of the term being evaluated.

The output of the elaboration process is a target language term that can be type checked with the rules from Figure 5.

If elaboration fails it can either be because the source level program describes a term that would be ill-typed when elaborated, or in some cases, elaboration fails because it cannot infer all the implicit parameters. This might happen if unification for the index language is undecidable, as is for example the case for contextual LF. In this case, annotations are needed when the term falls outside the strict pattern fragment where unification is decidable; this is rarely a problem in practice. For other index languages where unification is decidable, we do not expect such annotations to be necessary.

4. Description of elaboration

Elaboration of our source-language to our core target language is bi-directional and guided by the expected target type. Recall that we mark in the target type the arguments which are implicitly quantified (see $\Pi^i X:U.T$). This annotation is added when we elaborate a source type with free variables. If we check a source expression against $\Pi^i X:U.T$ we insert the appropriate λ -abstraction in our target. If we have synthesized the type $\Pi^i X:U.T$ for an expression, we insert *hole variables* for the omitted argument of type U . When we switch between synthesizing a type S for a given expression and checking an expression against an expected type T , we will rely on unification to make them equal. A key challenge is how to elaborate case-expressions where pattern matching a dependently typed expression of type τ against a pattern in a branch might refine the type τ . Our elaboration is parametric in the index domain, hence we keep our definitions of holes, instantiation of holes and unification abstract and only state their definitions and properties.

4.1 Elaboration of index objects

To elaborate a source expression, we insert holes for omitted index arguments and elaborate index objects which occur in it. We characterize holes with contextual objects as in [Pientka 2009]. Contextual objects encode the dependencies on the context that the hole might have. We hence make a few requirements about our index domain. We assume:

1. A function $\text{genHole} (?Y:\Delta.U)$ that generates a term standing for a hole of type U in the context Δ , i.e. its instantiation may refer to the index variables in Δ . If the index language is first-order, then we can characterize holes for example by meta-variables [Nanevski et al. 2008]. If our index language is higher-order, for example if we choose contextual LF as in Beluga, we characterize holes using meta²-variables as described in Boespflug and Pientka [2011]. As is common in these meta-variable calculi, holes are associated with a delayed substitution θ which is applied as soon as we know what Y stands for.
2. A typing judgment for guaranteeing that index objects with holes are well-typed:

$$\Theta; \Delta \vdash C:U \quad \text{Index object } C \text{ has index type } U \text{ in context } \Delta \\ \text{and all holes in } C \text{ are declared in } \Theta$$

where Θ stores the hole typing assumptions:

$$\text{Hole Context } \Theta ::= \cdot \mid \Theta, ?X:\Delta.U$$

3. Unification algorithm which finds the most general unifier for two index objects. In Beluga, we rely on the higher-order unification; more specifically, we solve eagerly terms which fall into the pattern fragment [Miller 1991; Dowek et al. 1996] and delay others [Abel and Pientka 2011]. A most general unifier exists if all unification constraints can be solved. Our elaboration relies on unifying computation-level types which in turn relies on unifying index-level terms; technically, we in fact rely on two unification judgments: one finding instantiations for holes in Θ , the other finding most general instantiations for index variables defined in Δ such that two index terms become equal. We use the first one during elaboration when unifying two computation-level types; the second one is used when computing the type refinement in branches.

$$\begin{array}{ll} \Theta; \Delta \vdash C_1 \doteq C_2 / \Theta'; \rho & \text{where: } \Theta' \vdash \rho: \Theta \\ \Delta \vdash C_1 \doteq C_2 / \Delta'; \theta & \text{where: } \Delta' \vdash \theta: \Delta \end{array}$$

where ρ describes the instantiation for holes in Θ . If unification succeeds, then we have $\llbracket \rho \rrbracket C_1 = \llbracket \rho \rrbracket C_2$ and $[\theta]C_1 = [\theta]C_2$ respectively.

4. Elaboration of index objects themselves. If the index language is simply typed, the elaboration has nothing to do; however, if as in Beluga, our index objects are objects described in the logical framework LF, then we need to elaborate them and infer omitted arguments following [Pientka 2013]. There are two related forms of elaboration for index objects we use:

$$\begin{array}{l} \Theta; \Delta \vdash c:U \quad \rightsquigarrow C/\Theta'; \Delta'; \rho \\ \Theta; \Delta \vdash \{c; \theta\}:U \rightsquigarrow C/\Theta'; \rho \end{array}$$

The first judgment reconstructs the index object c by checking it against U . We thread through a context Θ of holes and a context of index variables Δ , we have seen so far. The object c however may contain additional free index variables whose type we infer during elaboration. All variables occurring in C will be eventually declared with their corresponding type in Δ' . As we elaborate c , we may refine holes and add additional holes. ρ describes the mapping between Θ and Θ' , i.e. it records refinement of holes. Finally, we know that $\Delta' = \llbracket \rho \rrbracket \Delta, \Delta_0$, i.e. Δ' is an extension of Δ . We use the first judgment in elaborating patterns and type declarations in the signature.

The second judgment is similar to the first, but does not allow free index variables in c . We elaborate c together with a refinement substitution θ , which records refinements obtained from earlier branches. When we encounter an index variable, we look up what it is mapped to in θ and return it. Given a hole context Θ and an index variable context Δ , we elaborate an index term c against a given type U . The result is two fold: a context Θ' of holes is related to the original hole context Θ via the hole instantiation ρ . We use the second judgment to elaborate index objects embedded into target expressions.

4.2 Elaborating declarations

We begin our discussion of elaborating source programs in a top-down manner starting with declarations, the entry point of the algorithm. Types and kinds in declarations may contain free variables and there are two different tasks: we need to fill in omitted arguments, infer the type of free variables and abstract over the free variables and holes which are left over in the elaborated type/kind. We rely here on the fact that the index language provides a way of inferring the type of free variables.

To abstract over holes in a given type T , we employ a lifting operation: $\Delta \vdash \epsilon: \Theta$ which maps each hole to a fresh index variable.

$$\frac{}{\cdot \vdash \cdot: \cdot} \quad \frac{\Delta \vdash \epsilon: \Theta}{\Delta, X:U \vdash \epsilon, (.X)/X: \Theta, X: (.U)}$$

We require that holes are closed (written as $.U$ and $.X$ resp. where the context associated with a hole is empty); otherwise lifting fails². In other words, holes are not allowed to depend on some local meta-variables.

We use double brackets (i.e. $\llbracket \epsilon \rrbracket M$) to represent the application of the lifting substitutions and hole instantiation substitutions. We use this to distinguish them from regular substitutions such as the refinement substitutions in the target language.

Elaborating declarations requires three judgements. One for constants and one for kinds to be able to reconstruct inductive type

²In our implementation of elaboration in Beluga, we did not find this restriction to matter in practice.

$$\begin{array}{c}
\frac{\cdot; \cdot \vdash t \rightsquigarrow T/\Theta; \Delta; \cdot \quad \Delta_i \vdash \epsilon : \Theta}{\vdash \mathbf{c} : t \rightsquigarrow \Pi^i(\Delta_i, \llbracket \epsilon \rrbracket \Delta). \llbracket \epsilon \rrbracket T} \text{el-typ} \quad \frac{\cdot; \cdot \vdash k \rightsquigarrow K/\Theta; \Delta; \cdot \quad \Delta_i \vdash \epsilon : \Theta}{\vdash \mathbf{a} : k \rightsquigarrow \Pi^i(\Delta_i, \llbracket \epsilon \rrbracket \Delta). \llbracket \epsilon \rrbracket K} \text{el-kind} \\
\\
\frac{\cdot; \cdot \vdash t \rightsquigarrow T/\Theta; \Delta; \cdot \quad \Delta_i \vdash \epsilon : \Theta \quad \cdot; f : \Pi^i \Delta_i, \llbracket \epsilon \rrbracket \Delta. T \vdash \lambda e ; \cdot : \Pi^i(\Delta_i, \llbracket \epsilon \rrbracket \Delta). \llbracket \epsilon \rrbracket T \rightsquigarrow E/\cdot; \cdot}{\vdash \mathbf{rec} f : t = e \rightsquigarrow \mathbf{rec} f : \Pi^i(\Delta_i, \llbracket \epsilon \rrbracket \Delta). \llbracket \epsilon \rrbracket T = E} \text{el-rec}
\end{array}$$

Figure 6. Elaborating declarations

declarations, and one for recursive functions. These judgements are:

$$\begin{array}{ll}
\vdash \mathbf{c} : t & \rightsquigarrow T \\
\vdash \mathbf{a} : k & \rightsquigarrow K \\
\vdash \mathbf{rec} f : t = e & \rightsquigarrow \mathbf{rec} f : T = E
\end{array}$$

The elaboration of declarations succeeds when the result does not contain holes.

Figure 6 shows the rules for elaborating declarations. To elaborate a constant declarations $\mathbf{c} : t$ we elaborate the type t to a target type T where free index variables are listed in Δ and the remaining holes in T are described in Θ . We then lift all the holes in Θ to proper declarations in Δ_i via the lifting substitution ϵ . The final elaborated type of the constant \mathbf{c} is: $\Pi^i(\Delta_i, \llbracket \epsilon \rrbracket \Delta). \llbracket \epsilon \rrbracket T$. Note that both the free variables in the type t and the lifted holes described in Δ_i form the implicit arguments and are marked with Π^i . For example in the certifying evaluator, the type of the constructor Ex is reconstructed to:

$$\begin{array}{l}
\Pi^i T : [\text{tp}], M : [\text{term } T]. \\
\Pi^e N : [\text{value } T]. [\text{big-step } T \ M \ N] \rightarrow \text{Cert } [T] [M]
\end{array}$$

The elaboration of kinds follows the same principle.

To elaborate recursive function declarations, we first elaborate the type t abstracting over all the free variables and lifting the remaining holes to obtain $\Pi^i(\Delta_i, \llbracket \epsilon \rrbracket \Delta). \llbracket \epsilon \rrbracket T$. Second, we assume f of this type and elaborate the body e checking it against $\Pi^i(\Delta_i, \llbracket \epsilon \rrbracket \Delta). \llbracket \epsilon \rrbracket T$. We note that we always elaborate a source expression e together with a possible refinement substitution θ . In the beginning, θ will be empty. We describe elaboration of source expressions in the next section.

4.3 Elaborating source expressions

We elaborate source expressions bidirectionally. Expressions such as non-dependent functions and dependent functions are elaborated by checking the expression against a given type; expressions such as application and dependent application are elaborated to a corresponding target expression and at the same time synthesize the corresponding type. This approach seeks to propagate the typing information that we know in the checking rules, and in the synthesis phase, to take advantage of types that we can infer.

$$\begin{array}{ll}
\text{Synthesizing:} & \Theta; \Delta; \Gamma \vdash \lambda e ; \theta \rightsquigarrow E : T/\Theta'; \rho \\
\text{Checking:} & \Theta; \Delta; \Gamma \vdash \lambda e ; \theta : T \rightsquigarrow E \ / \Theta'; \rho
\end{array}$$

We first explain the judgment for elaborating a source expression e by checking it against T given holes in Θ , index variables Δ , and program variables Γ . Because of pattern matching, index variables in Δ may get refined to concrete index terms. Abusing slightly notation, we write θ for the map of free variables occurring in e to their refinements and consider a source expression e together with the refinement map θ , written as $\lambda e ; \theta$. The result of elaborating $\lambda e ; \theta$ is a target expression E , a new context of holes Θ' , and a hole instantiation ρ which instantiates holes in Θ , i.e. $\Theta' \vdash \rho : \Theta$. The result E has type $\llbracket \rho \rrbracket T$.

The result of elaboration in synthesis mode is similar; we return the target expression E together with its type T , a new context of holes Θ' and a hole instantiation ρ , s.t. $\Theta' \vdash \rho : \Theta$. The result is well-typed, i.e. E has type T .

We give the rules for elaborating source expressions in checking mode in Fig. 7 and in synthesis mode in Fig. 8. To elaborate a function (see rule **el-fn**) we simply elaborate the body extending the context Γ . There are two cases when we elaborate an expression of dependent function type. In the rule **el-mlam**, we elaborate a dependent function $\lambda X \Rightarrow e$ against $\Pi^e X : U. T$ by elaborating the body e extending the context Δ with the declaration $X : U$. In the rule **el-mlam-i**, we elaborate an expression e against $\Pi^i X : U. T$ by elaborating e against T extending the context Δ with the declaration $X : U$. The result of elaborating e is then wrapped in a dependent function.

When switching to synthesis mode, we elaborate $\lambda e ; \theta$ and obtain the corresponding target expression E and type T' together with an instantiation ρ for holes in Θ . We then unify the synthesized type T' and the expected type $\llbracket \rho \rrbracket T$ obtaining an instantiation ρ' and return the composition of the instantiation ρ and ρ' . When elaborating an index object $[c]$ (see rule **el-box**), we resort to elaborating c in our indexed language which we assume.

One of the key cases is the one for case-expressions. In the rule **el-case**, we elaborate the scrutinee synthesizing a type S ; we then elaborate the branches. Note that we verify that S is a closed type, i.e. it is not allowed to refer to holes. To put it differently, the type of the scrutinee must be fully known. This is done to keep a type refinement in the branches from influencing the type of the scrutinee. The practical impact of this restriction is difficult to quantify, however this seems to be the case for the programs we want to write as it is not a problem in any of the examples of the Beluga implementation. For a similar reason, we enforce that the type T , the overall type of the case-expression, is closed; were we to allow holes in T , we would need to reconcile the different instantiations found in different branches.

We omit the special case of pattern matching on index objects to save space and because it is a refinement on the **el-case** rule where we keep the scrutinee when we elaborate a branch. We then unify the scrutinee with the pattern in addition to unifying the type of the scrutinee with the type of the pattern. In our implementation in Beluga, we distinguish between case-expressions on computation-level expressions (which do not need to track the scrutinee and are described in the paper) and case-expressions on index objects (which do keep the scrutinee when elaborating branches).

When elaborating a constant, we look up its type T_c in the signature Σ and then insert holes for the arguments marked implicit in its type (see Fig. 8). Recall that all implicit arguments are quantified at the outside, i.e. $T_c = \Pi^i X_n : U_n. \dots \Pi^i X_1 : U_1. S$ where S does not contain any implicit dependent types Π^i . We generate for each implicit declaration $X_k : U_k$ a new hole which can depend on the currently available index variables Δ . When elaborating a variable, we look up its type in Γ and because the variable can correspond to a recursive function with implicit parameters we insert holes for the arguments marked as implicit as in the constant case.

$$\boxed{\Theta; \Delta; \Gamma \vdash \zeta e; \theta \rangle : T \rightsquigarrow E/\Theta'; \rho} \text{ Elaborate source } \zeta e; \theta \rangle \text{ to target expression } E \text{ checking against type } T$$

$$\frac{\Theta; \Delta \vdash \zeta c; \theta \rangle : U \rightsquigarrow C/\Theta'; \rho}{\Theta; \Delta; \Gamma \vdash \zeta [c]; \theta \rangle : [U] \rightsquigarrow [C]/\Theta'; \rho} \text{el-box} \quad \frac{\Theta; \Delta; \Gamma, x:T_1 \vdash \zeta e; \theta \rangle : T_2 \rightsquigarrow E/\Theta'; \rho}{\Theta; \Delta; \Gamma \vdash \zeta \mathbf{fn} x \Rightarrow e; \theta \rangle : T_1 \rightarrow T_2 \rightsquigarrow \mathbf{fn} x \Rightarrow E/\Theta'; \rho} \text{el-fn}$$

$$\frac{\Theta; \Delta, X:U; \Gamma \vdash \zeta e; \theta, X/X \rangle : T \rightsquigarrow E/\Theta'; \rho}{\Theta; \Delta; \Gamma \vdash \zeta e; \theta \rangle : \Pi^i X:U. T \rightsquigarrow \lambda X \Rightarrow E/\Theta'; \rho} \text{el-mlam-i} \quad \frac{\Theta; \Delta, X:U; \Gamma \vdash \zeta e; \theta, X/X \rangle : T \rightsquigarrow E/\Theta'; \rho}{\Theta; \Delta; \Gamma \vdash \zeta \lambda X \Rightarrow e; \theta \rangle : \Pi^e X:U. T \rightsquigarrow \lambda X \Rightarrow E/\Theta'; \rho} \text{el-mlam}$$

$$\frac{\Theta; \Delta; \Gamma \vdash \zeta e; \theta \rangle \rightsquigarrow E:S/\cdot; \rho \quad \llbracket \rho \rrbracket \Delta; \llbracket \rho \rrbracket \Gamma \vdash \zeta \vec{b}; \llbracket \rho \rrbracket \theta \rangle : S \rightarrow \llbracket \rho \rrbracket T \rightsquigarrow \vec{B}}{\Theta; \Delta; \Gamma \vdash \zeta \mathbf{case} e \mathbf{of} \vec{b}; \theta \rangle : T \rightsquigarrow \mathbf{case} E \mathbf{of} \vec{B}/\cdot; \rho} \text{el-case}$$

$$\frac{\Theta; \Delta; \Gamma \vdash \zeta e; \theta \rangle \rightsquigarrow E:T_1/\Theta_1; \rho \quad \Theta_1; \llbracket \rho \rrbracket \Delta \vdash T_1 \doteq \llbracket \rho \rrbracket T/\Theta_2; \rho'}{\Theta; \Delta; \Gamma \vdash \zeta e; \theta \rangle : T \rightsquigarrow \llbracket \rho' \rrbracket E/\Theta_2; \rho' \circ \rho} \text{el-syn}$$

Figure 7. Elaboration of Expressions (Checking Mode)

Elaboration of applications in the synthesis mode threads through the hole context and its instantiation, but is otherwise straightforward. In each of the application rules, we elaborate the first argument of the application obtaining a new hole context Θ_1 together with a hole instantiation ρ_1 . We then apply the hole instantiation ρ_1 to the context Δ and Γ and to the refinement substitution θ , before elaborating the second part.

4.3.1 Elaborating branches

We give the rules for elaborating branches in Fig. 9. Recall that a branch $pat \mapsto e$ consists of the pattern pat and the body e . We elaborate a branch under the refinement θ , because the body e may contain index variables declared earlier and which might have been refined in earlier branches.

Intuitively, to elaborate a branch, we need to elaborate the pattern and synthesize the type of index and pattern variables bound inside of it. In the dependently typed setting, pattern elaboration needs to do however a bit more work: we need to infer implicit arguments which were omitted by the programmer (e.g. the constructor \mathbf{Ex} takes the type of the expression, and the source of evaluation as implicit parameter $\mathbf{Ex} \ [T] \ [M] \ \dots$) and we need to establish how the synthesized type of the pattern refines the type of the scrutinee.

Moreover, there is a mismatch between the variables the body e may refer to (see rule $\mathbf{wf}\text{-branch}$ in Fig. 2) and the context in which the elaborated body E is meaningful (see rule $\mathbf{t}\text{-branch}$ in Fig. 5). While our source expression e possibly can refer to index variables declared prior, the elaborated body E is not allowed to refer to any index variables which were declared at the outside; those index variables are replaced by their corresponding refinements. To account for these additional refinements, we not only return an elaborated pattern $\Pi \Delta_r; \Gamma_r. Pat : \theta_r$ when elaborating a pattern pat (see rule $\mathbf{el}\text{-subst}$ in Fig. 9), but in addition return a map θ_e between source variables declared explicitly outside to their refinements.

Technically, elaborating a pattern is done in three steps (see rule $\mathbf{el}\text{-subst}$).

1. First, given pat we elaborate it to a target pattern Pat together with its type S_p synthesizing the type of index variables Δ_p and the type of pattern variables Γ_p together with holes (Θ_p) which denote omitted arguments. This is accomplished by the first premise of the rule $\mathbf{el}\text{-subst}$:

$$\cdot; \cdot \vdash pat \rightsquigarrow \Pi \Delta_p; \Gamma_p. Pat : S_1/\Theta_p; \cdot$$

Our pattern elaboration judgement (figure 10) threads through the hole context and the context of index variables, both of which are empty in the beginning. Because program variables occur linearly, we do not thread them through but simply combine program variable contexts when needed. The result of elaborating pat is a pattern Pat in our target language where Δ_p describes all index variables in Pat , Γ_p contains all program variables and Θ_p contains all holes, i.e. most general instantiations of omitted arguments. We describe pattern elaboration in detail in Section 4.3.2.

2. Second, we abstract over the hole variables in Θ_p by lifting all holes to fresh index variables from Δ'_p . This is accomplished by the second premise of the rule $\mathbf{el}\text{-subst}$.
3. Finally, we compute the refinement substitution θ_r which ensures that the type of the pattern $\llbracket \rho \rrbracket S_p$ is compatible with the type S of the scrutinee. We note that the type of the scrutinee could also force a refinement of holes in the pattern. This is accomplished by the judgment:

$$\Delta, (\Delta'_p, \llbracket \rho \rrbracket \Delta_p) \vdash \llbracket \rho \rrbracket S_1 \doteq T_1/\Delta_r; \theta_R \quad \theta_R = \theta_r, \theta_p$$

We note because θ_R maps index variables from $\Delta, (\Delta'_p, \llbracket \rho \rrbracket \Delta_p)$ to Δ_r , it contains two parts: θ_r provides refinements for variables Δ in the type of the scrutinee; θ_p provides possible refinements of the pattern forced by the scrutinee. This can happen, if the scrutinee's type is more specific than the type of the pattern.

4.3.2 Elaborating patterns

Pattern elaboration is bidirectional. The judgements for elaborating patterns by checking them against a given type and synthesizing their type are:

$$\begin{array}{ll}
\text{Synthesizing:} & \Theta; \Delta \vdash pat \rightsquigarrow \Pi \Delta'; \Gamma. Pat : T / \Theta'; \rho \\
\text{Checking:} & \Theta; \Delta \vdash pat : T \rightsquigarrow \Pi \Delta'; \Gamma. Pat / \Theta'; \rho
\end{array}$$

As mentioned earlier, we thread through a hole context Θ together with the hole substitution ρ that relates: $\Theta' \vdash \rho : \Theta$. Recall that as our examples show index-level variables in patterns need not to be linear and hence we accumulate index variables and thread them through as well. Program variables on the other hand must occur linearly, and we can simply combine them. The elaboration rules are presented in Figure 10. In synthesis mode, elaboration returns a reconstructed pattern Pat , a type T where Δ' describes the index variables in Pat and Γ' contains all program variables occurring in Pat . The hole context Θ' describes the most general instantiations for omitted arguments which have been inserted into

$$\boxed{\Theta; \Delta \vdash E : T \rightsquigarrow E' : T' / \Theta'} \quad \text{Apply } E \text{ to holes for representing omitted arguments based on } T \text{ obtaining a term } E' \text{ of type } T'$$

$$\frac{\text{genHole } (?Y:\Delta.U) = C \quad (\Theta, ?Y:\Delta.U); \Delta \vdash E [C] : [C/X]T \rightsquigarrow E':T' / \Theta'}{\Theta; \Delta \vdash E : \Pi^i X:U.T \rightsquigarrow E':T' / \Theta'} \quad \text{el-impl} \quad \frac{S \neq \Pi^i X:U.T}{\Theta; \Delta \vdash E:S \rightsquigarrow E:S / \Theta} \quad \text{el-impl-done}$$

$$\boxed{\Theta; \Delta; \Gamma \vdash \lambda e; \theta \rightsquigarrow E:T/\Theta'; \rho} \quad \text{Elaborate source } \lambda e; \theta \rightsquigarrow \text{ to target } E \text{ and synthesize type } T$$

$$\frac{\Gamma(x) = T \quad \Theta; \Delta; \Gamma \vdash x:T \rightsquigarrow E':T' / \Theta'}{\Theta; \Delta; \Gamma \vdash \lambda x; \theta \rightsquigarrow E':T' / \Theta'; \text{id}(\Theta')} \quad \text{el-var} \quad \frac{\Sigma(c) = T_c \quad \Theta; \Delta \vdash \mathbf{c} : T_c \rightsquigarrow E : T / \Theta'}{\Theta; \Delta; \Gamma \vdash \lambda \mathbf{c}; \theta \rightsquigarrow E : T / \Theta'; \text{id}(\Theta')} \quad \text{el-const}$$

$$\frac{\Theta; \Delta; \Gamma \vdash \lambda e_1; \theta \rightsquigarrow E_1:S \rightarrow T / \Theta_1; \rho_1 \quad \Theta_1; \llbracket \rho_1 \rrbracket \Delta; \llbracket \rho_1 \rrbracket \Gamma \vdash \lambda e_2; \llbracket \rho_1 \rrbracket \theta \rightsquigarrow \llbracket \rho_1 \rrbracket S \rightsquigarrow E_2 / \Theta_2; \rho_2}{\Theta; \Delta; \Gamma \vdash \lambda e_1 e_2; \theta \rightsquigarrow E_1 E_2 : \llbracket \rho_2 \rrbracket T / \Theta_2; \rho_2 \circ \rho_1} \quad \text{el-app}$$

$$\frac{\Theta; \Delta; \Gamma \vdash \lambda e; \theta \rightsquigarrow E_1:\Pi^e X:U.T/\Theta_1; \rho_1 \quad \Theta_1; \llbracket \rho_1 \rrbracket \Delta \vdash \lambda c; \llbracket \rho_1 \rrbracket \theta \rightsquigarrow U \rightsquigarrow C/\Theta_2; \rho_2}{\Theta; \Delta; \Gamma \vdash \lambda e[c]; \theta \rightsquigarrow E_1 [C]:[C/X](\llbracket \rho_2 \rrbracket T)/\Theta_2; \rho_2 \circ \rho_1} \quad \text{el-mapp}$$

$$\frac{\Theta; \Delta; \Gamma \vdash \lambda e; \theta \rightsquigarrow E:\Pi^e X:U.T/\Theta_1; \rho \quad \text{genHole } (?Y : (\llbracket \rho \rrbracket \Delta).U) = C}{\Theta; \Delta; \Gamma \vdash \lambda e_-; \theta \rightsquigarrow E [C] : [C/X]T / \Theta_1, ?Y:(\llbracket \rho_1 \rrbracket \Delta).U; \rho} \quad \text{el-mapp-underscore}$$

$$\frac{\Theta; \Delta \vdash \lambda t; \theta \rightsquigarrow T/\Theta_1; \rho_1 \quad \Theta_1; \llbracket \rho_1 \rrbracket \Delta; \llbracket \rho_1 \rrbracket \Delta \vdash \lambda e; \llbracket \rho_1 \rrbracket \theta \rightsquigarrow T \rightsquigarrow E/\Theta_2; \rho_2}{\Theta; \Delta; \Gamma \vdash \lambda e:t; \theta \rightsquigarrow (E:T):T/\Theta_2; \rho_2 \circ \rho_1} \quad \text{el-annotated}$$

Where $\text{id}(\Theta)$ returns the identity substitution for context Θ such as: $\Theta \vdash \text{id}(\Theta) : \Theta$

Figure 8. Elaborating of Expressions (Synthesizing Mode)

$$\boxed{\Delta; \Gamma \vdash \lambda b; \theta \rightsquigarrow S \rightarrow T \rightsquigarrow B} \quad \text{Elaborate source branch } \lambda b; \theta \rightsquigarrow \text{ to target branch } B$$

$$\frac{\Delta \vdash \text{pat} : S \rightsquigarrow \Pi \Delta_r; \Gamma_r. \text{Pat} : \theta_r \mid \theta_e \quad \cdot; \Delta_r; [\theta_r] \Gamma, \Gamma_r \vdash \lambda e; \theta_r \circ \theta, \theta_e \rightsquigarrow [\theta_r] T \rightsquigarrow E / \cdot}{\Delta; \Gamma \vdash \lambda \text{pat} \mapsto e; \theta \rightsquigarrow S \rightarrow T \rightsquigarrow \Pi \Delta_r; \Gamma_r. \text{Pat} : \theta_r \mapsto E} \quad \text{el-branch}$$

$$\boxed{\Delta \vdash \text{pat} : T \rightsquigarrow \Pi \Delta_r; \Gamma_r. \text{Pat} : \theta_r \mid \theta_e}$$

$$\frac{\cdot; \cdot \vdash \text{pat} \rightsquigarrow \Pi \Delta_p; \Gamma_p. \text{Pat} : S_p / \Theta_p; \cdot \quad \Delta'_p \vdash \rho : \Theta_p \quad \Delta, (\Delta'_p, \llbracket \rho \rrbracket \Delta_p) \vdash \llbracket \rho \rrbracket S_p \div S / \Delta_r; \theta_R}{\Delta \vdash \text{pat} : S \rightsquigarrow \Pi \Delta_r; [\theta_p] \llbracket \rho \rrbracket \Gamma_p. [\theta_p] \llbracket \rho \rrbracket \text{Pat} : \theta_r \mid \theta_e} \quad \text{el-subst}$$

where $\theta_R = \theta_r, \theta_p$ s.t. $\Delta_r \vdash \theta_p : (\Delta'_p, \llbracket \rho \rrbracket \Delta_p)$ and $\theta_p = \theta_i, \theta_e$ s.t. $\Delta_r \vdash \theta_i : \Delta'_p$

Figure 9. Branches and patterns

Pat. In checking mode, we elaborate *pat* given a type *T* to the target expression *Pat* and index variable context Δ' , pattern variable context Γ' and the hole context Θ' .

Pattern elaboration starts in synthesis mode, i.e. either elaborating an annotated pattern ($e : t$) (see rule **el-pann**) or a pattern $\mathbf{c} \text{ pat}$ (see rule **el-pcon**). To reconstruct patterns that start with a constructor we first look-up the constructor in the signature Σ to get its fully elaborated type T_c and then elaborate the arguments $\vec{\text{pat}}$ against T_c . Elaborating the spine of arguments is guided by the type T_c . If $T_c = \Pi^i X:U.T$, then we generate a new hole for the omitted argument of type U . If $T_c = T_1 \rightarrow T_2$, then we elaborate the first argument in the spine pat against T_1 and the remaining arguments $\vec{\text{pat}}$ against T_2 . If $T_c = \Pi^e X:U.T$, then we elaborate the first argument in the spine $[c]$ against U and the remaining arguments $\vec{\text{pat}}$ against $[C/X]T$. When the spine is empty, denoted by \cdot , we simply return the final type and check that constructor was

fully applied by ensuring that the type S we reconstruct against is either of index level type, i.e. $[U]$, or a recursive type, i.e. $\mathbf{a}[\vec{C}]$.

For synthesizing the patterns with a type annotation, first we elaborate the type t in an empty context using a judgement that returns the reconstructed type T , its holes and index variables (contexts Θ' and Δ'). Once we have the type we elaborate the pattern checking against the type T .

To be able to synthesize the type of pattern variables and return it, we check variables against a given type T during elaboration (see rule **el-pvar**). For index level objects, rule **el-pindex** we defer to the index level elaboration that the index domain provides³. Finally, when elaborating a pattern against a given type it is possible to switch to synthesis mode using rule **el-psyn**, where first we

³ Both, elaboration of pattern variables and of index objects can be generalized by for example generating a type skeleton in the rule **el-subst** given the scrutinee's type. This is in fact what is done in the implementation of Beluga.

$$\begin{array}{c}
\text{Pattern (synthesis mode)} \quad \boxed{\Theta; \Delta \vdash pat \rightsquigarrow \Pi\Delta'; \Gamma.Pat:T / \Theta'; \rho} \\
\\
\frac{\Sigma(c) = T \quad \Theta; \Delta \vdash \overrightarrow{pat} : T \rightsquigarrow \Pi\Delta'; \Gamma.\overrightarrow{Pat} \rangle S / \Theta'; \rho}{\Theta; \Delta \vdash \mathbf{c}\overrightarrow{pat} \rightsquigarrow \Pi\Delta'; \Gamma.\mathbf{c}\overrightarrow{Pat}:S / \Theta'; \rho} \text{el-pcon} \\
\\
\frac{\cdot; \cdot \vdash \{t; \cdot\} \rightsquigarrow T/\Theta'; \Delta'; \cdot \quad (\Theta, \Theta'); (\Delta, \Delta') \vdash pat : T \rightsquigarrow \Pi\Delta'; \Gamma.Pat / \Theta''; \rho'}{\Theta; \Delta \vdash (pat:t) \rightsquigarrow \Pi\Delta''; \Gamma.Pat:\llbracket \rho' \rrbracket T / \Theta''; \rho'} \text{el-pann} \\
\\
\text{Pattern (checking mode)} \quad \boxed{\Theta; \Delta \vdash pat : T \rightsquigarrow \Pi\Delta'; \Gamma.Pat / \Theta'; \rho} \\
\\
\frac{}{\Theta; \Delta \vdash x : T \rightsquigarrow \Pi\Delta; x:T.x / \Theta; \text{id}(\Theta)} \text{el-pvar} \quad \frac{\Theta; \Delta \vdash c : U \rightsquigarrow C/\Theta'; \Delta'; \rho}{\Theta; \Delta \vdash [c] : [U] \rightsquigarrow \Pi\Delta'; \cdot [C]/\Theta'; \rho} \text{el-pindex} \\
\\
\frac{\Theta; \Delta \vdash pat \rightsquigarrow \Pi\Delta'; \Gamma.Pat:S / \Theta'; \rho \quad \Theta'; \Delta' \vdash S \doteq \llbracket \rho \rrbracket T / \rho'; \Theta''}{\Theta; \Delta \vdash pat : T \rightsquigarrow \Pi\llbracket \rho' \rrbracket \Delta'; \llbracket \rho' \rrbracket \Gamma . \llbracket \rho \rrbracket Pat / \Theta''; \rho' \circ \rho} \text{el-psyn} \\
\\
\text{Pattern Spines} \quad \boxed{\Theta; \Delta \vdash \overrightarrow{pat} : T \rightsquigarrow \Pi\Delta'; \Gamma.\overrightarrow{Pat} \rangle S / \Theta'; \rho} \\
\\
\frac{\text{either } T = [U] \text{ or } T = \mathbf{a} \llbracket C \rrbracket}{\Theta; \Delta \vdash \cdot : T \rightsquigarrow \Pi\Delta; \cdot \cdot \rangle T / \Theta; \text{id}(\Theta)} \text{el-sp-empty} \\
\\
\frac{\Theta; \Delta \vdash pat : T_1 \rightsquigarrow \Pi\Delta'; \Gamma.Pat/\Theta'; \rho \quad \Theta'; \Delta' \vdash \overrightarrow{pat} : \llbracket \rho \rrbracket T_2 \rightsquigarrow \Pi\Delta''; \Gamma'.\overrightarrow{Pat} \rangle S / \Theta''; \rho'}{\Theta; \Delta \vdash pat \overrightarrow{pat} : T_1 \rightarrow T_2 \rightsquigarrow \Pi\Delta''; (\Gamma, \Gamma') . (\llbracket \rho' \rrbracket Pat) \overrightarrow{Pat} \rangle S / \Theta''; \rho' \circ \rho} \text{el-sp-cmp} \\
\\
\frac{\Theta; \Delta \vdash c : U \rightsquigarrow C/\Theta'; \Delta'; \rho \quad \Theta'; \Delta' \vdash \overrightarrow{pat} : [C/X]\llbracket \rho \rrbracket T \rightsquigarrow \Pi\Delta''; \Gamma.\overrightarrow{Pat} \rangle S / \Theta''; \rho'}{\Theta; \Delta \vdash [c] \overrightarrow{pat} : \Pi^e X:U.T \rightsquigarrow \Pi\Delta''; \Gamma . (\llbracket \rho' \rrbracket [C]) \overrightarrow{Pat} \rangle S / \Theta''; \rho' \circ \rho} \text{el-sp-explicit} \\
\\
\frac{\text{genHole } (?Y:\Delta.U) = C \quad \Theta; ?Y:\Delta.U; \Delta \vdash \overrightarrow{pat} : [C/X]T \rightsquigarrow \Pi\Delta'; \Gamma.\overrightarrow{Pat} \rangle S / \Theta'; \rho}{\Theta; \Delta \vdash \overrightarrow{pat} : \Pi^i X:U.T \rightsquigarrow \Pi\Delta'; \Gamma.(\llbracket \rho \rrbracket C) \overrightarrow{Pat} \rangle S / \Theta'; \rho} \text{el-sp-implicit}
\end{array}$$

Figure 10. Elaboration of patterns and pattern spines

elaborate the pattern synthesizing its type S and then we make sure that S unifies against the type T it should check against.

5. Soundness of elaboration

We establish soundness of our elaboration: if we start with a well-formed source expression, we obtain a well-typed target expression E which may still contain some holes and E is well-typed for any ground instantiation of these holes. In fact, our final result of elaborating a recursive function and branches must always return a closed expression.

Theorem 1 (Soundness).

1. If $\Theta; \Delta; \Gamma \vdash \{e; \theta\} : T \rightsquigarrow E/\Theta_1; \rho_1$ then for any grounding hole instantiation ρ_g s.t. $\cdot \vdash \rho_g : \Theta_1$ and $\rho_0 = \rho_g \circ \rho_1$, we have $\llbracket \rho_0 \rrbracket \Delta; \llbracket \rho_0 \rrbracket \Gamma \vdash \llbracket \rho_g \rrbracket E \Leftarrow \llbracket \rho_0 \rrbracket T$.
2. If $\Theta; \Delta; \Gamma \vdash \{e; \theta\} \rightsquigarrow E:T/\Theta_1; \rho_1$ then for any grounding hole instantiation ρ_g s.t. $\cdot \vdash \rho_g : \Theta_1$ and $\rho_0 = \rho_g \circ \rho_1$, we have $\llbracket \rho_0 \rrbracket \Delta; \llbracket \rho_0 \rrbracket \Gamma \vdash \llbracket \rho_g \rrbracket E \Rightarrow \llbracket \rho_g \rrbracket T$.
3. If $\Delta; \Gamma \vdash \{pat \mapsto e; \theta\} : S \rightarrow T \rightsquigarrow \Pi\Delta'; \Gamma'.Pat : \theta' \mapsto E$ then $\Delta; \Gamma \vdash \Pi\Delta'; \Gamma'.Pat : \theta' \mapsto E \Leftarrow S \rightarrow T$.

To establish soundness of elaboration of case-expressions and branches, we rely on pattern elaboration which abstracts over the

variables in patterns as well as over the holes which derive from most general instantiations inferred for omitted arguments. We abstract over these holes using a lifting substitution ϵ . In practice, we need a slightly more general lemma than the one stated below which takes into account the possibility that holes in Pat are further refined (see Appendix).

Lemma 2 (Pattern elaboration).

1. If $\Theta; \Delta \vdash pat \rightsquigarrow \Pi\Delta_1; \Gamma_1.Pat:T/\Theta_1; \rho_1$ and ϵ is a ground lifting substitution, such as $\Delta_i \vdash \epsilon:\Theta_1$ then $\Delta_i, \llbracket \epsilon \rrbracket \Delta_1; \llbracket \epsilon \rrbracket \Gamma_1 \vdash \llbracket \epsilon \rrbracket Pat \Leftarrow \llbracket \epsilon \rrbracket T$.
2. If $\Theta; \Delta \vdash pat : T \rightsquigarrow \Pi\Delta_1; \Gamma_1.Pat/\Theta_1; \rho_1$ and ϵ is a ground lifting substitution, such as $\Delta_i \vdash \epsilon:\Theta_1$ then $\Delta_i, \llbracket \epsilon \rrbracket \Delta_1; \llbracket \epsilon \rrbracket \Gamma_1 \vdash \llbracket \epsilon \rrbracket Pat \Leftarrow \llbracket \epsilon \rrbracket \llbracket \rho_1 \rrbracket T$.
3. If $\Theta; \Delta \vdash \overrightarrow{pat} : T \rightsquigarrow \Pi\Delta_1; \Gamma_1.\overrightarrow{Pat} \rangle S/\Theta_1; \rho_1$ and ϵ is a ground lifting substitution, such as $\Delta_i \vdash \epsilon:\Theta_1$ then $\Delta_i, \llbracket \epsilon \rrbracket \Delta_1; \llbracket \epsilon \rrbracket \Gamma_1 \vdash \llbracket \epsilon \rrbracket \overrightarrow{Pat} \Leftarrow \llbracket \epsilon \rrbracket \llbracket \rho_1 \rrbracket T \rangle \llbracket \epsilon \rrbracket S$.

6. Related work

Our language contains indexed families of types that are related to Zenger's work [Zenger 1997] and the Dependent ML (DML) [Xi 2007] and Applied Type System (ATS) [Xi 2004; Chen and Xi 2005]. The objective in these systems is: a program that is typable

in the extended indexed type system is already typable in ML. By essentially erasing all the type annotations necessary for verifying the given program is dependently typed, we obtain a simply typed ML-like program. In contrast, our language supports pattern matching on index objects. Our elaboration, in contrast to the one given in Xi [2007], inserts omitted arguments producing programs in a fully explicit dependently typed core language. This is different from DML-like systems which treat *all* index arguments as implicit and do not provide a way for programmers to manipulate and pattern match directly on index objects. Allowing users to explicitly access and match on index arguments changes the game substantially.

Elaboration from implicit to explicit syntax for dependently typed systems has first been mentioned by Pollack [1990] although no concrete algorithm to reconstruct omitted arguments was given. Luther [2001] refined these ideas as part of the TYPELab project. He describes an elaboration and reconstruction for the calculus of constructions without treating recursive functions and pattern matching. There is in fact little work on elaborating dependently-typed source language supporting recursion and pattern matching. For example, the Agda bi-directional type inference algorithm described in Norell [2007] concentrates on a core dependently typed calculus enriched with dependent pairs, but omits the rules for its extension with recursion and pattern matching⁴. Idris, a dependently typed language developed by Brady [2013] uses a different technique. Idris starts by adding holes for all the implicit variables and it tries to instantiate these holes using unification. However, the language uses internally a tactic based elaborator that is exposed to the user who can interactively fill the holes using tactics. He does not prove soundness of the elaboration, but conjectures that given a type correct program its elaboration followed by a reverse elaboration produces a matching source level program.

A notable example, is the work by Asperti et al. [2012] on describing a bi-directional elaboration algorithm for the Calculus of (Co)Inductive Constructions (CCIC) implemented in Matita. Their setting is very different from ours: CCIC is more powerful than our language since the language of recursive programs can occur in types and there is no distinction between the index language and the programming language itself. Moreover in Matita, we are only allowed to write total programs and all types must be positive. For these reasons their source and target language is more verbose than ours and refinement, i.e. the translation of the source to the target, is much more complex than our elaboration. The difference between our language and Matita particularly comes to light when writing case-expressions. In Matita as in Coq, the programmer might need to supply an invariant for the scrutinee and the overall type of the case expression as a type annotation. Each branch then is checked against the type given in the invariant. Sometimes, these annotations can be inferred by using higher-order unification to find the invariant. In contrast, our case-expressions require no type annotations and we refine each branch according to refinement imposed by the pattern in each branch. The refinement is computed with help from higher-order unification. This makes our source and target language more light-weight and closer to a standard simply typed functional language.

Finally, refinement in Matita may leave some holes in the final program which then can be refined further by the user using for example tactics. We support no such interaction; in fact, we fail, if holes are left-over and the programmer is asked to provide more information.

Agda, Matita and Coq require users to abstract over all variables occurring in a type and the user statically labels arguments the user can freely omit. To ease the requirement of declaring all variables

occurring in type, many of these systems such as Agda supports simply listing the variables occurring in a declaration without the type. This however can be brittle since it requires that the user chose the right order. Moreover, the user has the possibility to locally override the implicit arguments mechanism and provide instantiations for implicit arguments explicitly. This is in contrast to our approach where we guide elaboration using type annotations and omit arguments based on the free variables occurring in the declared type, similarly to Idris which abstracts and makes implicit all the free variables in types.

This work is also related to type inference for Generalized Algebraic Data Types (i.e: GADTs) such as [Schrijvers et al. 2009]. Here the authors describe an algorithm where they try to infer the types of programs with GADTs when the principal type can be inferred and requiring type annotations for the cases that lack a principal type or it can not be inferred. This is in contrast to our system which always requires a type annotation at the function level. On the other hand our system supports a richer variety of index languages (some index languages can be themselves dependently typed as with Contextual LF in Beluga). Moreover we support pattern matching on index terms, a feature that is critical to enable reasoning about objects from the index level. Having said that, the approach to GADTs from [Schrijvers et al. 2009] offers interesting ideas for future work, first making the type annotations optional for cases when they can be inferred, and providing a declarative type systems that helps the programmer understand when will the elaboration succeed to infer the types.

7. Conclusion and future work

In this paper we describe a surface language for writing dependently typed programs where we separate the language of types and index objects from the language of programs. Total programs in our language correspond to first-order inductive proofs over a specific index domain where we mediate between the logical aspects and the domain-specific parts using a box modality. Our programming language supports indexed data-types, dependent pattern matching and recursion. Programmers can leave index variables free when declaring the type of a constructor or recursive program as a way of stating that arguments for these free variables should be inferred by the type-directed elaboration. This offers a lightweight mechanism for writing compact programs which resemble their ML counterparts and information pertaining to index arguments can be omitted. In particular, our handling of case-expressions does not require programmers to specify the type invariants the patterns and their bodies must satisfy and case expressions can be nested due to the refinement substitutions that mediate between the context inside and outside a branch. Moreover, we seamlessly support nested pattern matching inside functions in our surface and core languages (as opposed to languages such as Agda or Idris where the former supports pattern matching lambdas that are elaborated as top-level functions and the latter only supports simply typed nested pattern matching).

To guide elaboration and type inference, we allow type annotations which indirectly refine the type of sub-expressions; type annotations in patterns are also convenient to name index variables which do not occur explicitly in a pattern.

We prove our elaboration sound, in the sense that if elaboration produces a fully explicit term, this term will be well-typed. Finally, our elaboration is implemented in Beluga, where we use as the index domain contextual LF, and has been shown practical (see for example the implementation of a type-preserving compiler [Belanger et al. 2013]). We believe our language presents an interesting point in the design space for dependently typed languages in general and sheds light into how to design and implement a depen-

⁴Norell [2007] contains extensive discussions on pattern matching and recursion, but the chapter on elaboration does not discuss them.

dently typed language where we have a separate index language, but still want to support pattern matching on these indices.

As future work, we would like to strengthen our soundness result by relating the source and elaborated expressions with a type directed equivalence. Furthermore, we intend to explore an appropriate notion of completeness of elaboration. This would provide stronger guarantees for programmers stating that all terms in the target language can be written as terms in the source language such that elaboration succeeds. Also, it would be interesting to add type inference for functions when such inference is decidable and to precisely characterize when annotations can be omitted (again, related to the notion of completeness). Moreover, we would like to explore more powerful type-systems for the computational language, such as polymorphism. Polymorphism would improve how useful the language is, but it is largely an orthogonal feature that would not impact much what is discussed in this paper.

Acknowledgments

The authors would like to thank the anonymous reviewers for their feedback and constructive criticism. It really helped improve this work.

References

- A. Abel and B. Pientka. Higher-order dynamic pattern unification for dependent types and records. In L. Ong, editor, *10th International Conference on Typed Lambda Calculi and Applications (TLCA'11)*, Lecture Notes in Computer Science (LNCS 6690), pages 10–26. Springer, 2011.
- A. Asperti, W. Ricciotti, C. S. Coen, and E. Tassi. A bi-directional refinement algorithm for the calculus of (co)inductive constructions. *Logical Methods in Computer Science*, 8:1–49, 2012.
- O. S. Belanger, S. Monnier, and B. Pientka. Programming type-safe transformations using higher-order abstract syntax. In G. Gonthier and M. Norrish, editors, *Third International Conference on Certified Programs and Proofs (CPP'13)*, Lecture Notes in Computer Science (LCNS 8307), pages 243–258. Springer, 2013.
- M. Boespflug and B. Pientka. Multi-level contextual modal type theory. In G. Nadathur and H. Geuvers, editors, *6th International Workshop on Logical Frameworks and Meta-languages: Theory and Practice (LFMTP'11)*, Electronic Proceedings in Theoretical Computer Science (EPTCS), 2011.
- E. Brady. Idris, a general-purpose dependently typed programming language: Design and implementation. *Journal of Functional Programming*, 23:552–593, 9 2013.
- A. Cave and B. Pientka. Programming with binders and indexed data-types. In *39th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'12)*, pages 413–424. ACM Press, 2012.
- C. Chen and H. Xi. Combining programming with theorem proving. In O. Danvy and B. C. Pierce, editors, *10th International Conference on Functional Programming*, pages 66–77, 2005.
- G. Dowek, T. Hardin, C. Kirchner, and F. Pfenning. Unification via explicit substitutions: The case of higher-order patterns. In M. Maher, editor, *Joint International Conference and Symposium on Logic Programming*, pages 259–273. MIT Press, Sept. 1996.
- J. Dunfield and B. Pientka. Case analysis of higher-order data. In *International Workshop on Logical Frameworks and Meta-Languages: Theory and Practice (LFMTP'08)*, volume 228 of *Electronic Notes in Theoretical Computer Science (ENTCS)*, pages 69–84. Elsevier, June 2009.
- R. Harper, F. Honsell, and G. Plotkin. A framework for defining logics. *Journal of the ACM*, 40(1):143–184, January 1993.
- L. Jia, J. A. Vaughan, A. Mazurak, J. Zhao, L. Zarko, J. Schorr, and S. Zdancewic. Aura: A programming language for authorization and audit. In *13th ACM SIGPLAN International Conference on Functional Programming*, pages 27–38. ACM, 2008.
- M. Luther. More on implicit syntax. In R. Gore, A. Leitsch, and T. Nipkow, editors, *First International Joint Conference on Automated Reasoning (IJCAR'01)*, Lecture Notes in Artificial Intelligence (LNAI) 2083, pages 386–400. Springer, 2001.
- D. Miller. Unification of simply typed lambda-terms as logic programming. In *8th International Logic Programming Conference*, pages 255–269. MIT Press, 1991.
- A. Nanevski, F. Pfenning, and B. Pientka. Contextual modal type theory. *ACM Transactions on Computational Logic*, 9(3):1–49, 2008.
- U. Norell. *Towards a practical programming language based on dependent type theory*. PhD thesis, Department of Computer Science and Engineering, Chalmers University of Technology, Sept. 2007. Technical Report 33D.
- F. Pfenning. Elf: A language for logic definition and verified meta-programming. In *Fourth Annual Symposium on Logic in Computer Science*, pages 313–322, Pacific Grove, California, June 1989. IEEE Computer Society Press.
- F. Pfenning and C. Elliott. Higher-order abstract syntax. In *ACM SIGPLAN Symposium on Language Design and Implementation (PLDI'88)*, pages 199–208, June 1988.
- F. Pfenning and C. Schürmann. System description: Twelf — a meta-logical framework for deductive systems. In H. Ganzinger, editor, *16th International Conference on Automated Deduction (CADE-16)*, Lecture Notes in Artificial Intelligence (LNAI 1632), pages 202–206. Springer, 1999.
- B. Pientka. A type-theoretic foundation for programming with higher-order abstract syntax and first-class substitutions. In *35th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'08)*, pages 371–382. ACM Press, 2008.
- B. Pientka. Higher-order term indexing using substitution trees. *ACM Transactions on Computational Logic*, 11(1):1–40, 2009.
- B. Pientka. An insider's look at LF type reconstruction: Everything you (n)ever wanted to know. *Journal of Functional Programming*, 1(1–37), 2013.
- B. Pientka and J. Dunfield. Beluga: a framework for programming and reasoning with deductive systems (System Description). In J. Giesl and R. Haehnle, editors, *5th International Joint Conference on Automated Reasoning (IJCAR'10)*, Lecture Notes in Artificial Intelligence (LNAI 6173), pages 15–21. Springer-Verlag, 2010.
- B. Pientka, S. S. Ruan, and A. Abel. Structural recursion over contextual objects. Technical report, School of Computer Science, McGill, January 2014.
- B. C. Pierce and D. N. Turner. Local type inference. *ACM Transaction on Programming Languages and Systems*, 22(1):1–44, Jan. 2000.
- R. Pollack. Implicit syntax. Informal Proceedings of First Workshop on Logical Frameworks, Antibes, 1990.
- T. Schrijvers, S. Peyton Jones, M. Sulzmann, and D. Vytiniotis. Complete and decidable type inference for gadts. In *14th ACM SIGPLAN International Conference on Functional Programming*, ICFP '09, pages 341–352. ACM, 2009.
- K. Watkins, I. Cervesato, F. Pfenning, and D. Walker. A concurrent logical framework I: Judgments and properties. Technical Report CMU-CS-02-101, Department of Computer Science, Carnegie Mellon University, 2002.
- H. Xi. Applied type system. In *TYPES 2003*, volume 3085 of *Lecture Notes in Computer Science*, pages 394–408. Springer, 2004.
- H. Xi. Dependent ml an approach to practical programming with dependent types. *Journal of Functional Programming*, 17:215–286, 3 2007.
- H. Xi and F. Pfenning. Dependent types in practical programming. In *26th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'99)*, pages 214–227. ACM Press, 1999.
- C. Zenger. Indexed types. *Theoretical Computer Science*, 187(1-2):147–165, 1997.